

AutoGluon.Multimodal 0.6.0

Installation

AutoGluon ([GitHub](#)) requires `pip > 1.4` (upgrade by `pip install -U pip`). [More installation options](#). AutoGluon supports Python 3.7 to 3.9. Installation is available for Linux, MacOS, and Windows.

```
| pip install autogluon
```

Classification & Regression

MultiModalPredictor finetunes foundation models for solving classification and regression problems with image, text, and tabular features. Here, we use a simplified version [petfinder_for_tutorial](#) from the [PetFinder dataset](#). MultiModalPredictor automatically analyzes the columns in the input dataframe to detect categorical, numerical, text, and images (stored as paths).

```
| import pandas as pd
| train_data = pd.read_csv(
|     'petfinder_for_tutorial/train.csv', index_col=0)
| test_data = pd.read_csv(
|     'petfinder_for_tutorial/test.csv', index_col=0)
```

To train the model, just call `fit()`. We also support customization ([docs](#)).

```
| from autogluon.multimodal import MultiModalPredictor
| predictor = MultiModalPredictor(
|     problem_type="classification",
|     label="AdoptionSpeed"
| )
| predictor.fit(train_data)
| predictor.fit_summary() # Summarize the model
```

To evaluate on the validation data and run inference

```
| predictor.evaluate(test_data) # Evaluation
| predictor.predict(test_data) # Inference
```

Named-Entity Recognition

MultiModalPredictor supports named-entity recognition. We use [MIT movies corpus](#) to demonstrate the usage, which can be downloaded from [train.csv](#) and [test.csv](#).

```
| from autogluon.multimodal import MultiModalPredictor
| predictor = MultiModalPredictor(
|     problem_type="ner", label="entity_annotations"
| )
| # Train model
| predictor.fit("train.csv")
| # Evaluation
| predictor.evaluate("test.csv")
|
| # Inference
| sentence = "Game of Thrones is an American fantasy "
|           "drama television series created"
|           "by David Benioff"
| predictions = predictor.predict({
|     'text_snippet': [sentence]
| })
```

Object Detection

To use MultiModalPredictor for object detection, please first install additional dependencies by

```
| mim install mcv-full
| pip install mmdet
| pip install pycocotools-windows # if windows users
```

MultiModalPredictor supports common object detection data formats such as [VOC](#) and [COCO](#). Here we use the dataset [tiny_motorbike_coco](#) to demonstrate how to use MultiModalPredictor. The predictor natively supports json files in the COCO-format.

```
| train_path = "./Annotations/trainval_cocoformat.json"
| test_path = "./Annotations/test_cocoformat.json"
|
| from autogluon.multimodal import MultiModalPredictor
| model = "faster_rcnn_r50_fpn_2x_coco"
| predictor = MultiModalPredictor(
|     problem_type="object_detection",
|     # Set the backbone. This is optional.
|     hyperparameters={
|         "model.mmdet_image.checkpoint_name": model,
|     },
|     sample_data_path=train_path,
| )
```

To train an object detector,

```
| predictor.fit(train_path)
```

More options for the `fit` method ([docs](#)):

```
| predictor.fit(train_path,
|     # Limit the training time, in second
|     time_limit=600,
|     # Use a separate dataset to tune models.
|     tuning_data=val_data
| )
```

Once trained, we can evaluate it on the test set to obtain metrics,

```
| predictor.evaluate(test_path)
```

We can also predict on any image and visualize the detected bounding boxes with its confidence scores,

```
| pred = predictor.predict({"image": [test_image]})
```



Matching

MultiModalPredictor implements a flexible twin-tower architecture that can solve text-text, image-image, and text-image matching problems ([docs](#)). Example of extracting embeddings for semantic text matching:

```
| pip install ir_datasets # Install dataset package
| from autogluon.multimodal import MultiModalPredictor
| from autogluon.multimodal import utils
| import ir_datasets
| import pandas as pd
| dataset = ir_datasets.load("beir/fiqa/dev")
| docs_df = pd.DataFrame(dataset.docs_iter()) \
|     .set_index("doc_id")
|
| model = "sentence-transformers/all-MiniLM-L6-v2"
| predictor = MultiModalPredictor(
|     problem_type="text_similarity",
|     hyperparameters={
|         "model.hf_text.checkpoint_name": model,
|     }
| )
| doc_embedding = predictor.extract_embedding(docs_df)
| q_embedding = predictor.extract_embedding([
|     "what happened when the dot com bubble burst?"
| ])
| similarity = utils.compute_semantic_similarity(
|     q_embedding, doc_embedding
| )
| # Get the most relevant document
| docs_df['text'].iloc[similarity.argmax().numpy()]
```

In addition, you can finetune the matching model via relevance data. Here, we demonstrate the use-case via the [Flickr30K](#) image-text matching dataset preprocessed in the dataframe format: [flickr30k.zip](#).

```
| import pandas as pd
| train_data = pd.read_csv("train.csv", index_col=0)
| tdata = pd.read_csv("test.csv", index_col=0)
```

To finetune model, just specify the “query” and “response” keys when creating predictor and pick “image_text_similarity” as problem type.

```
| from autogluon.multimodal import MultiModalPredictor
| predictor = MultiModalPredictor(
|     query="caption",
|     response="image",
|     problem_type="image_text_similarity",
| )
| predictor.fit(train_data, time_limit=180)
| e_i = predictor.extract_embedding(tdata["image"])
| e_t = predictor.extract_embedding(tdata["caption"])
```

- MultiModalPredictor also supports features like knowledge distillation ([docs](#)), efficient finetuning ([docs](#)), and HPO ([docs](#)).
- For deployment, check [Tutorial](#).
- For other use-cases, check [TabularPredictor](#) and [TimeSeriesPredictor](#).
- Check the [latest version of this cheat sheet](#).
- Any questions? [Ask here](#)
- Like what you see? Consider [starring AutoGluon on GitHub](#) and [following us on twitter](#) to get notified of the latest updates!